



Neural Programmer-Interpreters

Scott Reed & Nando de Freitas

Presented by Benjamín Farías V.



Contents

- 1. Context**
2. Related Work
3. Model
4. Experiments
5. Conclusions
6. Personal Criticism



Context - Machine Learning

Machine Learning Process



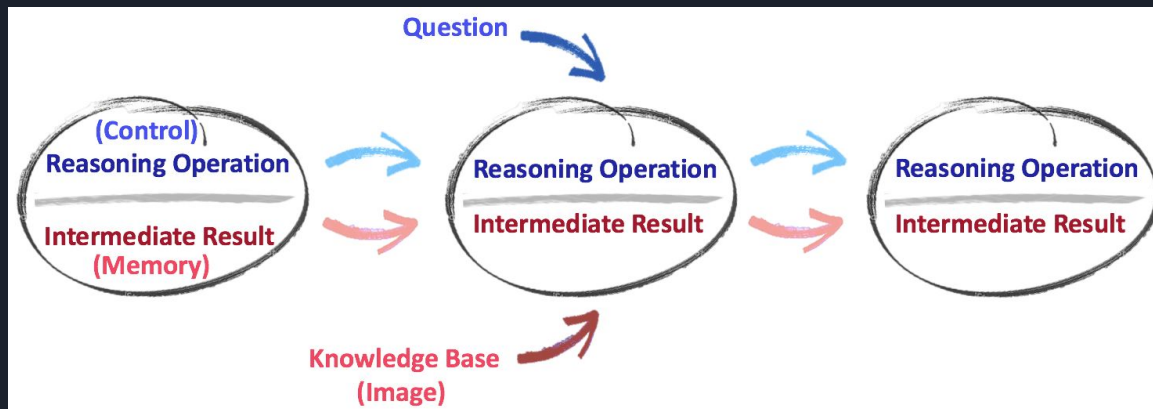
Pros:

- Good results
- Extensively studied

Cons:

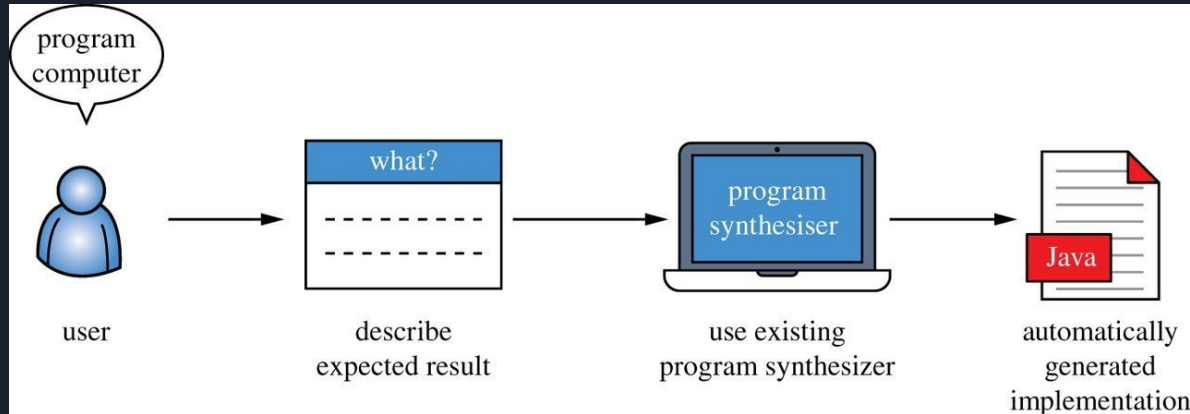
- Learns superficial patterns
- Requires a ton of data

Context - Machine Reasoning



- Learns logical rules from data
- Closer to human learning
- Requires less data

Context - Program Learning



- Networks that can deduce and learn programs
- Machines could create their own programs!
- **Future:** Human Level AI?




Contents

1. Context
- 2. Related Work**
3. Model
4. Experiments
5. Conclusions
6. Personal Criticism




Related Work - RNNs

Neural Turing Machines



Learn & Execute Simple Programs

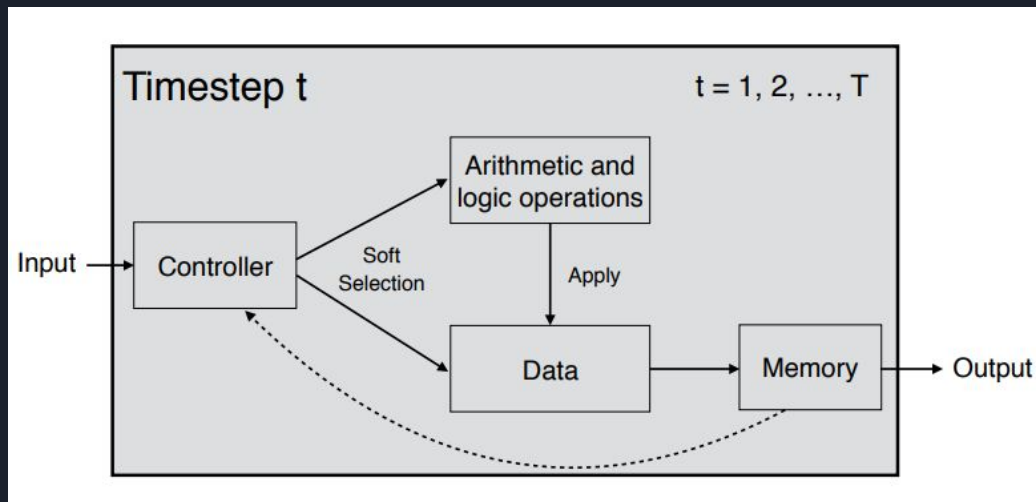
Pointer Networks



Output Space Depends on Input

Related Work - Program Induction

Neural Programmer



RNN + Controller + Operation + Memory



Related Work - Program Induction

Curriculum Learning

Humans and animals learn much better when the examples are not randomly presented but organized in a meaningful order which illustrates gradually more concepts, and gradually more complex ones. . . . and call them “curriculum learning”.

Bengio et al. (2009)

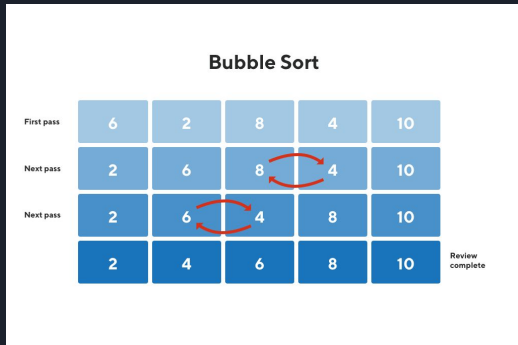


Contents

1. Context
2. Related Work
- 3. Model**
4. Experiments
5. Conclusions
6. Personal Criticism




Model - Neural Programmer-Interpreter (NPI)



[0.32, 0.77, 0.67, ..., 0.42]


- Learn to **represent and interpret** programs
- **Programmer:** Learns new program representations
- **Interpreter:** Executes learned programs over more complex tasks

Model - NPI Core



- **Multi-Layer LSTM** network
- Acts as a program router
- Decides which program to call next


Model - NPI Core



Input Components

- **State:** Environment observation + program arguments
- **Program:** Current program embedding

Model - NPI Core




Program

State

Input Components

- The **State** is obtained from a domain-specific encoder
- The **Program** is obtained from a memory module


Model - NPI Core



Hidden Component

- Receives the last hidden state (**$h-1$**)
- Computes the feed-forward step

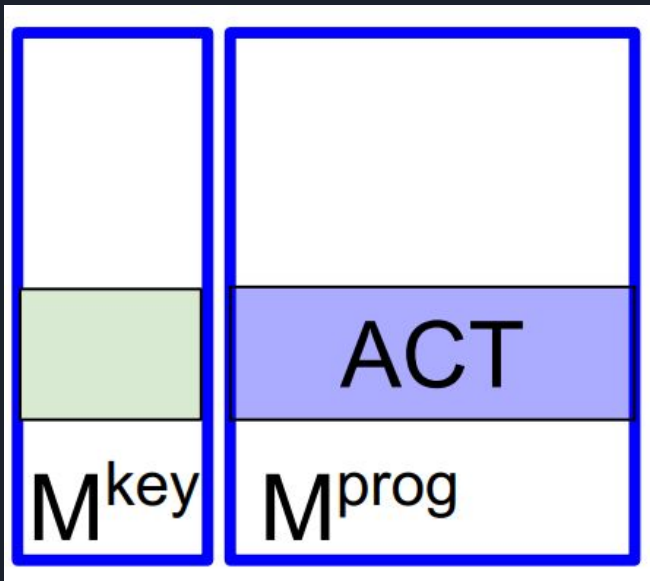
Model - NPI Core



Output Components

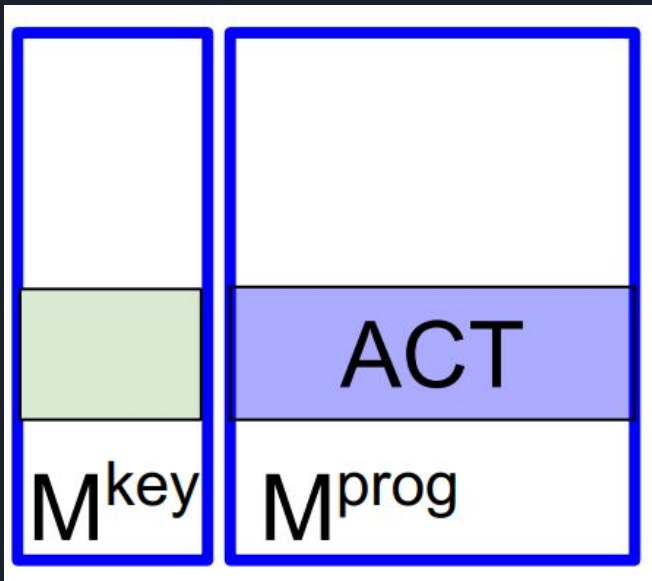
- **Key:** Lookup key embedding for next program
- **End:** Probability of returning
- **Arg:** Arguments for next program

Model - NPI Memory



- Global memory, has two components
- Each row in the **Key** component corresponds to the same row in the **Prog** component

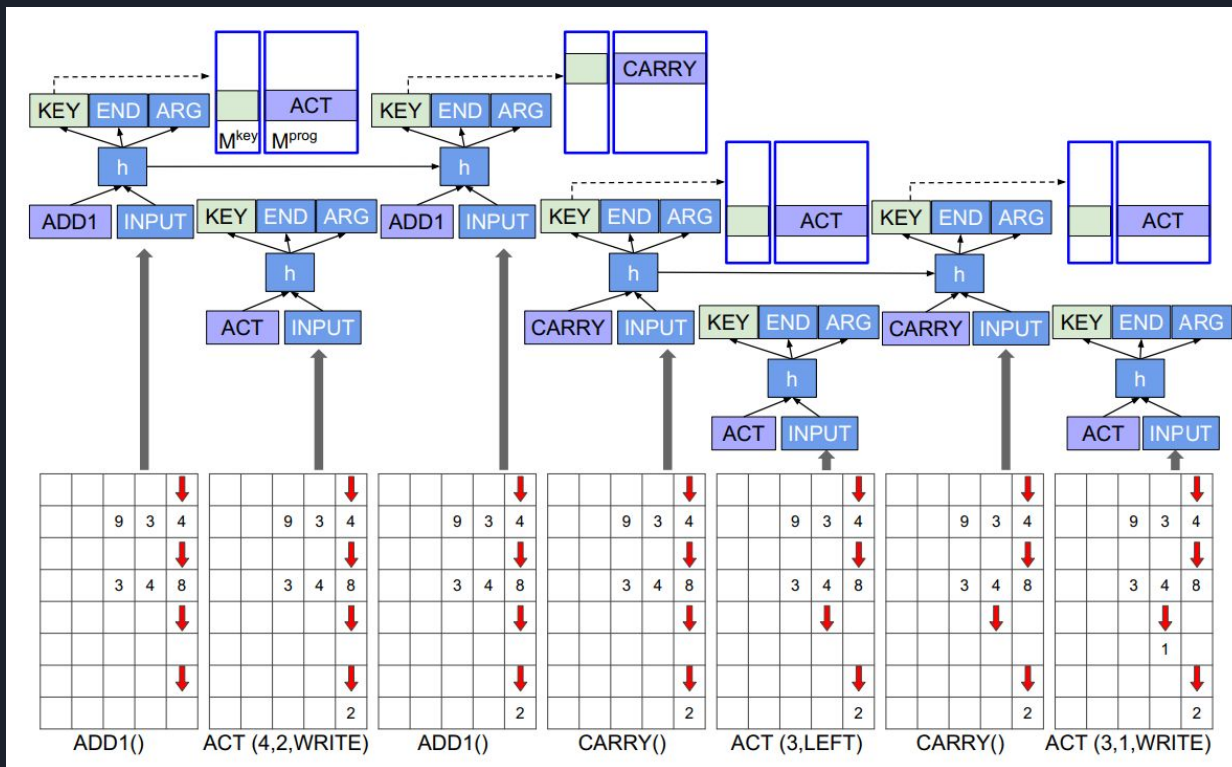
Model - NPI Memory



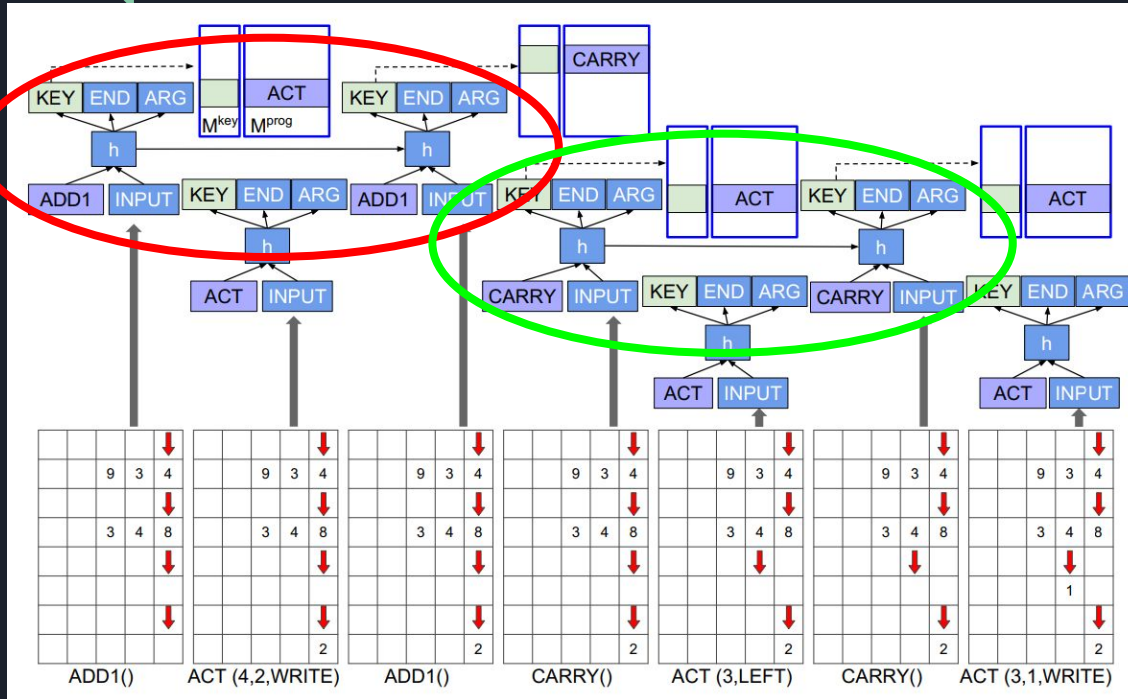
Memory Components

- **Key:** Stores all program keys
- **Prog:** Stores all program embeddings

Model - NPI Network (Addition)



Model - NPI Network (Addition)

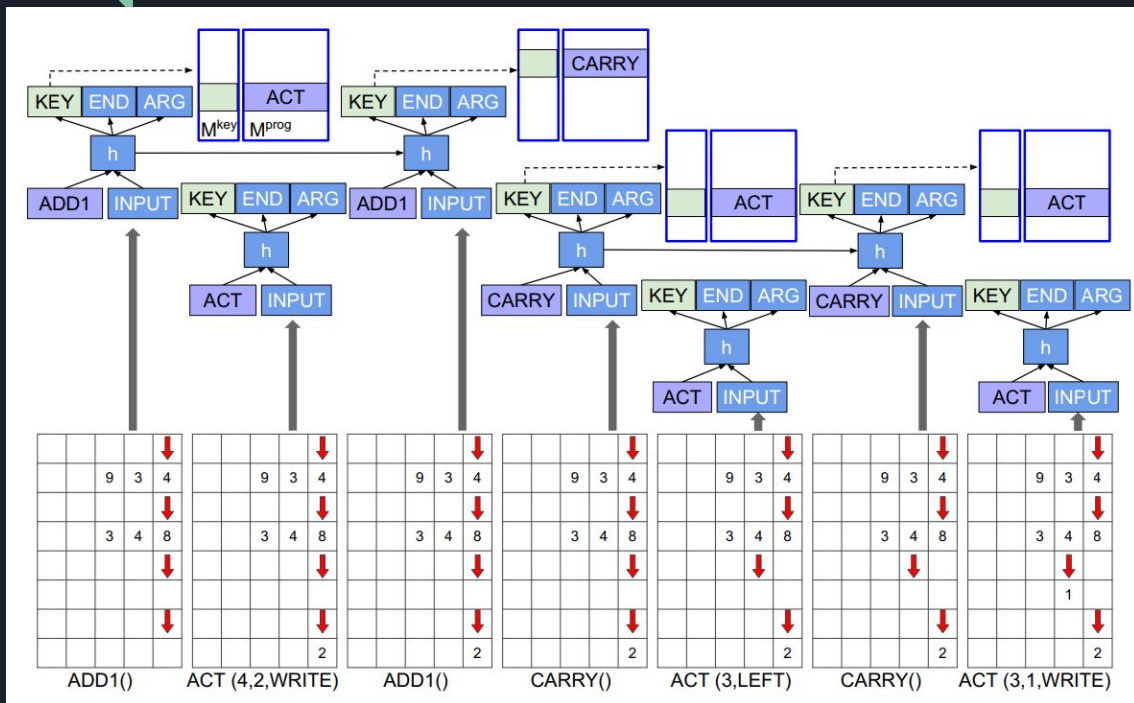


- Each program has one **NPI Core Network**
- **All NPI Cores** share the same weights
- Works like a **call stack**

ADD1()

CARRY()

Model - NPI Network (Addition)



[ADD1]

[ADD1, ACT]

[ADD1]

[ADD1, CARRY]

[ADD1, CARRY, ACT]

[ADD1, CARRY]

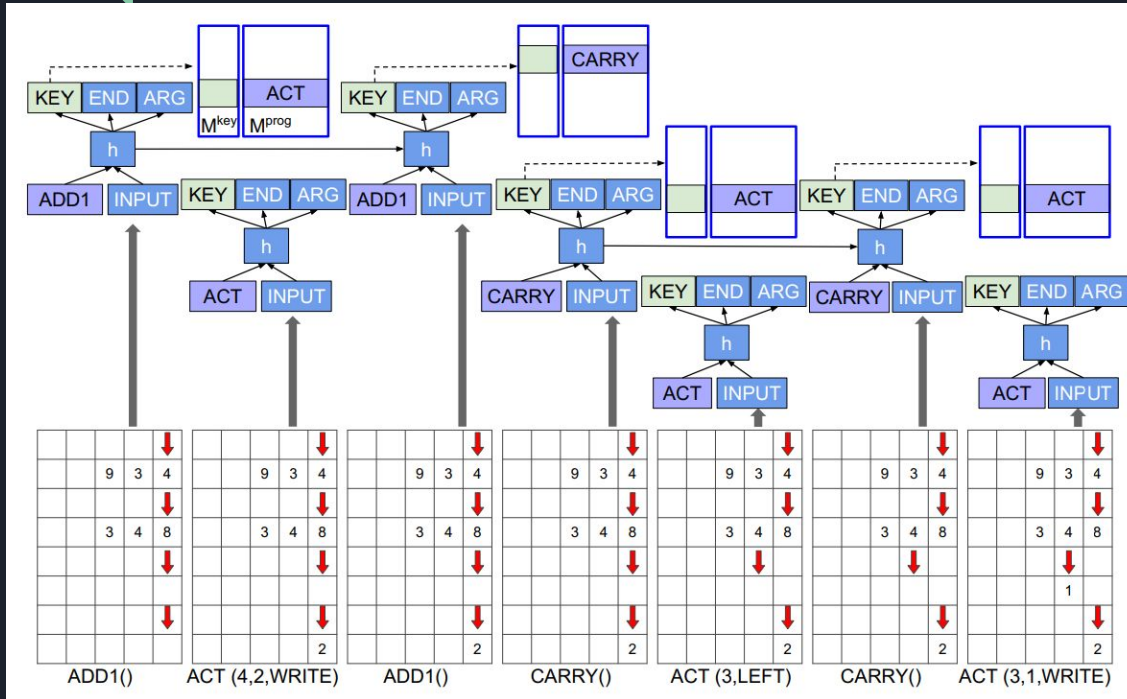
[ADD1, CARRY, ACT]

[ADD1, CARRY]

[ADD1]

[]

Model - NPI Network (Addition)



- The memory is shared
- The output is a sequence of **actions**

1. ACT (4, 2, WRITE)
2. ACT (3, LEFT)
3. ACT (3, 1, WRITE)



Model - NPI Training

```
ADD1  
  WRITE OUT 2  
LSHIFT  
  PTR INP1 LEFT  
  PTR INP2 LEFT  
  PTR CARRY LEFT  
  PTR OUT LEFT
```

- Use execution traces for real programs
- Predict the next program to be called
- Apply **curriculum learning** to focus on programs that the model is failing at



Contents

1. Context
2. Related Work
3. Model
- 4. Experiments**
5. Conclusions
6. Personal Criticism



Experiments - Addition

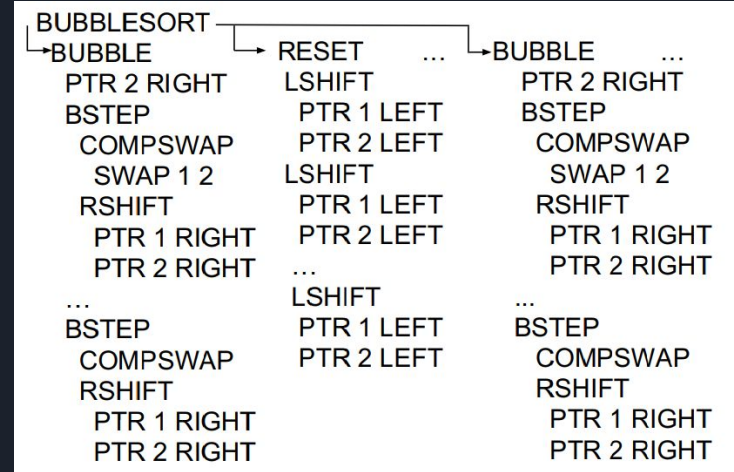

input 1	0	0	0	9	6
input 2	0	0	1	2	5
carry	0	0	1	1	1
output	0	0	0	2	1



ADD	ADD1	ADD1	ADD1
WRITE OUT 1	WRITE OUT 2	WRITE OUT 2	WRITE OUT 2
CARRY	CARRY	CARRY	LSHIFT
PTR CARRY LEFT	PTR CARRY LEFT	PTR CARRY LEFT	PTR INP1 LEFT
WRITE CARRY 1	WRITE CARRY 1	WRITE CARRY 1	PTR INP2 LEFT
PTR CARRY RIGHT	PTR CARRY RIGHT	PTR CARRY RIGHT	PTR CARRY LEFT
LSHIFT	LSHIFT	LSHIFT	PTR OUT LEFT
PTR INP1 LEFT	PTR INP1 LEFT	PTR INP1 LEFT	
PTR INP2 LEFT	PTR INP2 LEFT	PTR INP2 LEFT	
PTR CARRY LEFT	PTR CARRY LEFT	PTR CARRY LEFT	
PTR OUT LEFT	PTR OUT LEFT	PTR OUT LEFT	

- Addition of two base-10 numbers using a **scratch pad**
- The model can **move** the pointers and **write** numbers
- **Testing:** Addition for numbers with more digits

Experiments - Sorting



- Array sorting with **Bubble Sort** on a **scratch pad**
- The model can **move** the pointers and **swap** elements
- **Testing:** Sorting of longer arrays

Experiments - Canonicalize 3D Models




```
GOTO 1 2
  HGOTO
    RGOTO
      ACT(RIGHT)
      ACT(RIGHT)
      ACT(RIGHT)
    VGOTO
      DGOTO
        ACT(DOWN)
        ACT(DOWN)
```

- Move the **camera** to the target view by looking at the **image**
- The model can only see the **current rendering** of the car
- **Testing:** New car models and different positions

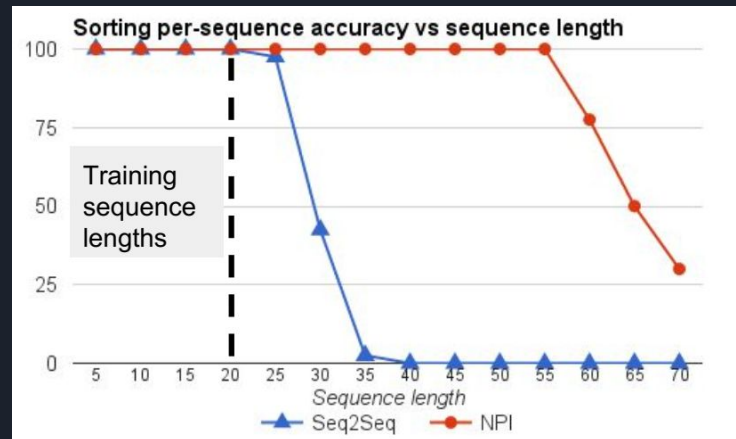
Experiments - Results (Sorting)

Accuracy VS # Training Examples



NPI learns at a way **faster rate** compared to a **Seq2Seq LSTM**

Accuracy VS Sequence Length



NPI generalizes to **longer sequences** compared to a **Seq2Seq LSTM**



Experiments - Results (Multitasking)

Task	Single	Multi	+ Max
Addition	100.0	97.0	97.0
Sorting	100.0	100.0	100.0
Canon. seen car	89.5	91.4	91.4
Canon. unseen	88.7	89.9	89.9
Maximum	-	-	100.0

- Single-Task models perform really well
- The **Multi-Task** model is comparable to **all** single-task models!
- **MAX** can be learned **without affecting performance** on previous tasks!



Contents

1. Context
2. Related Work
3. Model
4. Experiments
- 5. Conclusions**
6. Personal Criticism





Conclusions

- The **NPI** can learn programs from very dissimilar environments
- **Strong generalization** in comparison to **Seq2Seq LSTMs**
- A trained **NPI** with a fixed core can continue to learn without forgetting



Contents

1. Context
2. Related Work
3. Model
4. Experiments
5. Conclusions
- 6. Personal Criticism**





Personal Criticism

Good:

- Interesting approach to program learning by using composition
- Experiments show good generalization capabilities with little data

Bad:

- The model architecture is hard to understand from their explanation
- Training requires already having an implementation for each program



Bibliography

- [1] Reed, S., De Freitas, N. (2016). *Neural Programmer-Interpreters*.
- [2] Graves, A., Wayne, G., Danihelka, I. (2014). *Neural Turing Machines*.
- [3] Vinyals, O., Jaitly, N., Fortunato, M. (2015). *Pointer Networks*.
- [4] Bengio, Y., Louradour, J., Collobert, R., Weston, J. (2009). *Curriculum Learning*.
- [5] Neelakantan, A., Le, Q., Sutskever, I. (2016). *Neural Programmer: Inducing Latent Programs With Gradient Descent*.